

Modular Toolkit for Data Processing (MDP)

This document is also available online: <http://mdp-toolkit.sourceforge.net/tutorial.html>

Tutorial

Author: Pietro Berkes and Tiziano Zito

Homepage: <http://mdp-toolkit.sourceforge.net>

Copyright: This document has been placed in the public domain.

Version: 2.2

This is a guide to basic and some more advanced features of the MDP library. Besides the present tutorial, you can learn more about MDP by using the standard Python tools. All MDP nodes have doc-strings, the public attributes and methods have telling names: All information about a node can be obtained using the `help` and `dir` functions within the Python interpreter. In addition to that, an automatically generated [API](#) is available.

Note

Code snippets throughout the script will be denoted by:

```
>>> print "Hello world!"  
Hello world!
```

To run the following code examples don't forget to import `mdp` in your Python session with:

```
>>> import mdp
```

You'll find all the code of this tutorial within the `demo` directory in the MDP installation path.

Contents

[Tutorial](#)

[Introduction](#)

[Quick Start](#)

[Nodes](#)

[Node Instantiation](#)

[Node Training](#)

[Node Execution](#)

[Node Inversion](#)

[Writing your own nodes: subclassing Node](#)

[Flows](#)

- Flow instantiation, training and execution
- Flow inversion
- Flows are container type objects
- Crash recovery
- Iterators
- Checkpoints
- A real life example (Logistic maps)
- Another real life example (Growing neural gas)
- Node List
- Contributed Nodes
- Additional utilities
 - Graph module
- To Do
- Contributors

Introduction

Modular toolkit for Data Processing (MDP) is a data processing framework written in Python.

From the user's perspective, MDP consists of a collection of trainable supervised and unsupervised algorithms or other data processing units (nodes) that can be combined into data processing flows. Given a sequence of input data, MDP takes care of successively training or executing all nodes in the flow. This structure allows to specify complex algorithms as a sequence of simpler data processing steps in a natural way. Training can be performed using small chunks of input data, so that the use of very large data sets becomes possible while reducing the memory requirements. Memory usage can also be minimized by defining the internals of the nodes to be single precision.

The base of readily available algorithms includes Principal Component Analysis, two flavors of Independent Component Analysis, Slow Feature Analysis, Gaussian Classifiers, Growing Neural Gas, Fisher Discriminant Analysis, and Factor Analysis. The full list of implemented nodes can be found in the [Node List](#) section.

From the developer's perspective, MDP is a framework to make the implementation of new algorithms easier. The basic class 'Node' takes care of tedious tasks like numerical type and dimensionality checking, leaving the developer free to concentrate on the implementation of the training and execution phases. The node then automatically integrates with the rest of the library and can be used in a flow together with other nodes. A node can have multiple training phases and even an undetermined number of phases. This allows for example the implementation of algorithms that need to collect some statistics on the whole input before proceeding with the actual training, or others that need to iterate over a training phase until a convergence criterion is satisfied. The ability to train each phase using chunks of input data is maintained if the chunks are generated with iterators. Moreover, crash recovery is optionally available: in case of failure, the current state of the flow is saved for later inspection.

MDP has been written in the context of theoretical research in neuroscience, but it has been designed to be helpful in any context where trainable data processing algorithms are used. Its simplicity on the user side together with the reusability of the implemented nodes make it also a valid educational tool.

As its user base is steadily increasing, MDP appears as a good candidate for becoming a common repository of user-supplied, freely available, Python implemented data processing algorithms.

Quick Start

Using MDP is as easy as:

```

>>> import mdp
>>> # perform pca on some data x
...
>>> y = mdp.pca(x)
>>> # perform ica on some data x using single precision
...
>>> y = mdp.fastica(x, dtype='float32')

```

A complete list of all short-cut functions like `pca` or `fastica` can be obtained as follows:

```

>>> dir(mdp.helper_funcs)
['__builtins__', '__doc__', '__file__', '__name__',
'cubica', 'factor_analysis', 'fastica', 'get_eta',
'isfa', 'mdp', 'pca', 'sfa', 'sfa2', 'whitening']

```

MDP is, of course, much more than this: it allows to combine different algorithms and other data processing elements (nodes) into data processing sequences (flows). Moreover, it provides a framework that makes the implementation of new algorithms easy and intuitive.

MDP requires the numerical Python extensions [numpy](#) or [scipy](#). In its namespace MDP offers references to the main modules `numpy` or `scipy`, and the subpackages `linalg`, `random`, and `fft` as `mdp.numx`, `mdp.numx_linalg`, `mdp.numx_rand`, and `mdp.numx_fft`. This is done to possibly support additional numerical extensions in the future. At import time MDP will select `scipy` if available, otherwise `numpy` will be loaded. You can force the use of a numerical extension by setting the environment variable `MDPNUMX=numpy` or `MDPNUMX=scipy`.

Nodes

A node is the basic unit in MDP and it represents a data processing element, like for example a learning algorithm, a filter, a visualization step, etc. Each node can have one or more training phases, during which the internal structures are learned from training data (e.g. the weights of a neural network are adapted or the covariance matrix is estimated) and an execution phase, where new data can be processed forwards (by processing the data through the node) or backwards (by applying the inverse of the transformation computed by the node if defined). The `Node` class is designed to make the implementation of new algorithms easy and intuitive, for example by setting automatically input and output dimension and by casting the data to match the numerical type (e.g. float or double) of the internal structures. `Node` was designed to be applied to arbitrarily long sets of data: the internal structures can be updated incrementally by sending chunks of the input data (this is equivalent to online learning if the chunks consists of single observations, or to batch learning if the whole data is sent in a single chunk). A `Node` can be copied or saved using the corresponding `copy` and `save` methods.

Node Instantiation

Nodes can be obtained by creating an instance of the `Node` class. Each node is characterized by an input dimension, that corresponds to the dimensionality of the input vectors, an output dimension, and a `dtype`, which determines the numerical type of the internal structures and of the output signal. These three attributes are inherited from the input data if left unspecified. Input dimension and `dtype` can usually be specified when an instance of the node class is created. The constructor of each node class can require other task-specific arguments. The full documentation is available in the doc-string of the node's class.

Some examples of node instantiation:

- Create a node that performs Principal Component Analysis (PCA) whose input dimension and `dtype` are inherited from the input data during training. Output dimensions default to input dimensions.

```
>>> pcanode1 = mdp.nodes.PCANode()
>>> pcanode1
PCANode(input_dim=None, output_dim=None, dtype=None)
```

- Setting `output_dim = 10` means that the node will keep only the first 10 principal components of the input.

```
>>> pcanode2 = mdp.nodes.PCANode(output_dim = 10)
>>> pcanode2
PCANode(input_dim=None, output_dim=10, dtype=None)
```

The output dimensionality can also be specified in terms of the explained variance. If we want to keep the number of principal components which can account for 80% of the input variance, we set:

```
>>> pcanode3 = mdp.nodes.PCANode(output_dim = 0.8)
>>> pcanode3.desired_variance
0.80000000000000004
```

- If `dtype` is set to `float32` (32-bit float), the input data is cast to single precision when received and the internal structures are also stored as `float32`. `dtype` influences the memory space necessary for a node and the precision with which the computations are performed.

```
>>> pcanode4 = mdp.nodes.PCANode(dtype = 'float32')
>>> pcanode4
PCANode(input_dim=None, output_dim=None, dtype='float32')
```

You can obtain a list of the numerical types supported by a node looking at its `supported_dtypes` property:

```
>>> pcanode4.supported_dtypes
[dtype('float32'), dtype('float64')]
```

This method returns a list of `numpy.dtype` objects (see the `numpy` documentation for more details).

- A `PolynomialExpansionNode` expands its input in the space of polynomials of a given degree by computing all monomials up to the specified degree. Its constructor needs as first argument the degree of the polynomials space (3 in this case).

```
>>> expnode = mdp.nodes.PolynomialExpansionNode(3)
```

Node Training

Some nodes need to be trained to perform their task. This can be done during a training phases by calling the `train` method. MDP supports both supervised and unsupervised training, and algorithms with multiple training phases.

Some examples of node training:

- Create some random data and update the internal structures (i.e. mean and covariance matrix) of the `PCANode`:

```
>>> x = mdp.numx_rand.random((100, 25)) # 25 variables, 100 observations
>>> pcanode1.train(x)
```

At this point the input dimension and the `dtype` have been inherited from `x`:

```
>>> pcanode1
PCANode(input_dim=25, output_dim=None, dtype='float64')
```

- We can train our node with more than one chunk of data. This is especially useful when the input data is too long to be stored in memory or when it has to be created on-the-fly. (See also the [Iterators](#) section):

```
>>> for i in range(100):
...     x = mdp.numx_rand.random((100, 25))
...     pcanode1.train(x)
>>>
```

- Some nodes don't need to or cannot be trained:

```
>>> expnode.is_trainable()
False
```

Trying to train them anyway would raise an `IsNotTrainableException`.

- The training phase ends when the `stop_training`, `execute`, `inverse`, and possibly some other node-specific methods are called. For example we can stop the training of `pcanode1` (at this point the principal components are computed):

```
>>> pcanode1.stop_training()
```

- If the `PCANode` was declared to have a number of output components dependent on the input variance to be explained, we can check after training the number of output components and the actually explained variance:

```
>>> pcanode3.train(x)
>>> pcanode3.stop_training()
>>> pcanode3.output_dim
16
>>> pcanode3.explained_variance
0.85261144755506446
```

It is now possible to access the trained internal data. In general, a list of the interesting internal attributes can be found in the class documentation.

```
>>> avg = pcanode1.avg # mean of the input data
>>> v = pcanode1.get_projmatrix() # projection matrix
```

- Some nodes, namely the one corresponding to supervised algorithms, e.g. Fisher Discriminant Analysis (FDA), may need some labels or other supervised signals to be passed during training. Detailed information about the signature of the `train` method can be read in its doc-string.

```
>>> fdanode = mdp.nodes.FDANode()
>>> for label in ['a', 'b', 'c']:
...     x = mdp.numx_rand.random((100, 25))
...     fdanode.train(x, label)
>>>
```

- A node could also require multiple training phases. For example, the training of `fdanode` is not complete yet, since it has two training phases. We need to stop the first phase and train the second:

```

>>> fdanode.stop_training()
>>> for label in ['a', 'b', 'c']:
...     x = mdp.numx_rand.random((100, 25))
...     fdanode.train(x, label)
>>>

```

The easiest way to train multiple phase nodes is using [Flows](#) , which automatically handle multiple phases.

Node Execution

After the training phase it is possible to execute the node:

- The input data is projected on the principal components learned in the training phase:

```

>>> x = mdp.numx_rand.random((100, 25))
>>> y_pca = pcancode1.execute(x)

```

- Calling a node instance is equivalent to executing it:

```

>>> y_pca = pcancode1(x)

```

- The input data is expanded in the space of polynomials of degree 3:

```

>>> x = mdp.numx_rand.random((100, 5))
>>> y_exp = expnode(x)

```

- The input data is projected to the directions learned by FDA:

```

>>> x = mdp.numx_rand.random((100, 25))
>>> y_fda = fdanode(x)

```

- Some nodes may allow for optional arguments in the `execute` method, as always the complete information is given in the doc-string.

Node Inversion

If the operation computed by the node is invertible, it is possible to compute the inverse transformation:

- Given the output data, compute the inverse projection to the input space for the PCA node:

```

>>> pcancode1.is_invertible()
True
>>> x = pcancode1.inverse(y_pca)

```

- The expansion node is not invertible:

```

>>> expnode.is_invertible()
False

```

Trying to compute the inverse would raise an `IsNotInvertibleException`.

Writing your own nodes: subclassing Node

MDP tries to make it easy to write new data processing elements that fit with the existing elements. To expand the MDP library of implemented nodes with your own nodes you can subclass the Node class, overriding some of the methods according to your needs.

It is recommended to refer to the `numpy` or `scipy` numerical extensions through the MDP aliases `mdp.numx`, `mdp.numx_linalg`, `mdp.numx_fft`, and `mdp.numx_rand` when writing Node subclasses. This shall ensure that your nodes can be used without modifications should MDP support alternative numerical extensions in the future.

We'll illustrate this with some toy examples.

- We start by defining a node that multiplies its input by 2.

Define the class as a subclass of Node:

```
>>> class TimesTwoNode(mdp.Node):
```

This node cannot be trained. To specify this, one has to overwrite the `is_trainable` method to return False:

```
...     def is_trainable(self): return False
```

Execute only needs to multiply x by 2

```
...     def _execute(self, x):
...         return 2*x
```

Note that the `execute` method, which should never be overwritten and which is inherited from the Node parent class, will perform some tests, for example to make sure that `x` has the right rank, dimensionality and casts it to have the right `dtype`. After that the user-supplied `_execute` method is called. Each subclass has to handle the `dtype` defined by the user or inherited by the input data, and make sure that internal structures are stored consistently. To help with this the Node base class has a method called `_refcast(array, dtype)` that casts an array only when its `dtype` is different from the requested one.

The inverse of the multiplication by 2 is of course the division by 2:

```
...     def _inverse(self, y):
...         return y/2
...
>>>
```

Test the new node:

```
>>> node = TimesTwoNode(dtype = 'int32')
>>> x = mdp.numx.array([[1.0, 2.0, 3.0]])
>>> y = node(x)
>>> print x, '* 2 = ', y
[ [ 1.  2.  3.] ] * 2 = [ [ 2  4  6] ]
>>> print y, '/ 2 =', node.inverse(y)
[ [ 2  4  6] ] / 2 = [ [ 1  2  3] ]
```

- We then define a node that raises the input to the power specified in the initializer:

```
>>> class PowerNode(mdp.Node):
```

We redefine the `init` method to take the power as first argument. In general one should always give the possibility to set the `dtype` and the input dimensions. The default value is `None`, which means that the exact value is going to be inherited from the input data:

```
...     def __init__(self, power, input_dim=None, dtype=None):
```

Initialize the parent class:

```
...         super(PowerNode, self).__init__(input_dim=input_dim, dtype=dtype)
```

Store the power:

```
...         self.power = power
```

PowerNode is not trainable...

```
...     def is_trainable(self): return False
```

... nor invertible:

```
...     def is_invertible(self): return False
```

It is possible to overwrite the function `_get_supported_dtypes` to return a list of `dtype` supported by the node:

```
...     def _get_supported_dtypes(self):
...         return ['float32', 'float64']
```

The supported types can be specified in any format allowed by `numpy.dtype`. The interface method `get_supported_dtypes` converts them and sets the property `supported_dtypes`, which is a list of `dtype` objects.

The `_execute` method:

```
...     def _execute(self, x):
...         return self._refcast(x**self.power)
...
>>>
```

Test the new node

```
>>> node = PowerNode(3)
>>> x = mdp.numx.array([[1.0, 2.0, 3.0]])
>>> y = node.execute(x)
>>> print x, '**', node.power, '=', node(x)
[ [ 1.  2.  3.] ] ** 3 = [ [ 1.  8. 27.] ]
```

- We now define a node that needs to be trained. The `MeanFreeNode` computes the mean of its training data and subtracts it from the input during execution:

```
>>> class MeanFreeNode(mdp.Node):
...     def __init__(self, input_dim=None, dtype=None):
...         super(MeanFreeNode, self).__init__(input_dim=input_dim,
...                                             dtype=dtype)
```

We store the mean of the input data in an attribute. We initialize it to `None` since we still don't know how large is an input vector:

```
...         self.avg = None
```

Same for the number of training points:

```
...         self.tlen = 0
```


The subclass only needs to overwrite the `_train` method, which will be called by the parent `train` after some testing and casting has been done:

```
...     def _train(self, x):
...         # Initialize the mean vector with the right
...         # size and dtype if necessary:
...         if self.avg is None:
...             self.avg = mdp.numx.zeros(self.input_dim,
...                                       dtype=self.dtype)
```

Update the mean with the sum of the new data:

```
...         self.avg += mdp.numx.sum(x, axis=0)
```

Count the number of points processed:

```
...         self.tlen += x.shape[0]
```

Note that `train` method can have further arguments, which might be useful to implement algorithms that require supervised learning. For example, if you want to define a node that performs some form of classification you can define a `_train(self, data, labels)` method. The parent `train` checks `data` and takes care to pass the `labels` on (cf. for example `mdp.nodes.FDANode`).

The `_stop_training` function is called by the parent `stop_training` method when the training phase is over. We divide the sum of the training data by the number of training vectors to obtain the mean:

```
...     def _stop_training(self):
...         self.avg /= self.tlen
...         if self.output_dim is None:
...             self.output_dim = self.input_dim
```

Note that we `input_dim` are set automatically by the `train` method, and we want to ensure that the node has `output_dim` set after training. For nodes that do not need training, the setting is performed automatically upon execution. The `_execute` and `_inverse` methods:

```
...     def _execute(self, x):
...         return x - self.avg
...     def _inverse(self, y):
...         return y + self.avg
...
>>>
```

Test the new node:

```
>>> node = MeanFreeNode()
>>> x = mdp.numx_rand.random((10,4))
>>> node.train(x)
>>> y = node.execute(x)
>>> print 'Mean of y (should be zero): ', mdp.numx.mean(y, 0)
Mean of y (should be zero): [ 0.00000000e+00  2.22044605e-17
-2.22044605e-17  1.11022302e-17]
```

- It is also possible to define nodes with multiple training phases. In such a case, calling the `train` and `stop_training` functions multiple times is going to execute successive training phases (this kind of node is much easier to train using [Flows](#)). Here we'll define a node that returns a meanfree, unit variance signal. We define two training phases: first we compute the mean of the signal and next we sum the squared, meanfree input to compute the standard deviation (of course it is possible to solve this problem in one single step - remeber this is just a toy example).

```
>>> class UnitVarianceNode(mdp.Node):
...     def __init__(self, input_dim=None, dtype=None):
...         super(UnitVarianceNode, self).__init__(input_dim=input_dim,
...                                                 dtype=dtype)
...         self.avg = None # average
...         self.std = None # standard deviation
...         self.tlen = 0
```

The training sequence is defined by the user-supplied function `_get_train_seq`, that returns a list of tuples, one for each training phase. The tuples contain references to the training and stop-training functions of each of them. The default output of this function is `[(_train, _stop_training)]`, which explains the standard behavior illustrated above. We overwrite the function to return the list of our training functions:

```
...     def _get_train_seq(self):
...         return [(self._train_mean, self._stop_mean),
...                 (self._train_std, self._stop_std)]
```

Next we define the training functions. The first phase is identical to the one in the previous example:

```
...     def _train_mean(self, x):
...         if self.avg is None:
...             self.avg = mdp.numx.zeros(self.input_dim,
...                                       dtype=self.dtype)
...             self.avg += mdp.numx.sum(x, 0)
...             self.tlen += x.shape[0]
...     def _stop_mean(self):
...         self.avg /= self.tlen
```

The second one is only marginally different and does not require many explanations:

```
...     def _train_std(self, x):
...         if self.std is None:
...             self.tlen = 0
...             self.std = mdp.numx.zeros(self.input_dim,
...                                       dtype=self.dtype)
...             self.std += mdp.numx.sum((x - self.avg)**2., 0)
...             self.tlen += x.shape[0]
...     def _stop_std(self):
...         # compute the standard deviation
...         self.std = mdp.numx.sqrt(self.std/(self.tlen-1))
```

The `_execute` and `_inverse` methods are not surprising, either:

```
...     def _execute(self, x):
...         return (x - self.avg)/self.std
...     def _inverse(self, y):
...         return y*self.std + self.avg
>>>
```

Test the new node:

```
>>> node = UnitVarianceNode()
>>> x = mdp.numx_rand.random((10,4))
>>> # loop over phases
```

```

... for phase in range(2):
...     node.train(x)
...     node.stop_training()
...
...
>>> # execute
... y = node.execute(x)
>>> print 'Standard deviation of y (should be one): ', mdp.numx.std(y, axis=0)
Standard deviation of y (should be one): [ 1.  1.  1.  1.]

```

- In our last example we'll define a node that returns two copies of its input. The output is going to have twice as many dimensions.

```

>>> class TwiceNode(mdp.Node):
...     def is_trainable(self): return False
...     def is_invertible(self): return False

```

When `Node` inherits the input dimension, output dimension, and `dtype` from the input data, it calls the methods `set_input_dim`, `set_output_dim`, and `set_dtype`. Those are the setters for `input_dim`, `output_dim` and `dtype`, which are Python [properties](#). If a subclass needs to change the default behaviour, the internal methods `_set_input_dim`, `_set_output_dim` and `_set_dtype` can be overwritten. The property setter will call the internal method after some basic testing and internal settings. The private methods `_set_input_dim`, `_set_output_dim` and `_set_dtype` are responsible for setting the private attributes `_input_dim`, `_output_dim`, and `_dtype` that contain the actual value.

Here we overwrite `_set_input_dim` to automatically set the output dimension to be twice the input one, and `_set_output_dim` to raise an exception, since the output dimension should not be set explicitly.

```

...     def _set_input_dim(self, n):
...         self._input_dim = n
...         self._output_dim = 2*n
...     def _set_output_dim(self, n):
...         raise mdp.NodeException, "Output dim can not be set explicitly!"

```

The `_execute` method:

```

...     def _execute(self, x):
...         return mdp.numx.concatenate((x, x), 1)
...
>>>

```

Test the new node

```

>>> node = TwiceNode()
>>> x = mdp.numx.zeros((5,2))
>>> x
array([[0, 0],
       [0, 0],
       [0, 0],
       [0, 0],
       [0, 0]])
>>> node.execute(x)
array([[0, 0, 0, 0],

```

```
[0, 0, 0, 0],
[0, 0, 0, 0],
[0, 0, 0, 0],
[0, 0, 0, 0]])
```

Flows

A flow consists in an acyclic graph of nodes (currently only node sequences are implemented). The data is sent to an input node and is successively processed by the following nodes on the graph. The general flow implementation automatizes the training, execution, and inverse execution (if defined) of the whole graph. Training can be supervised and can consist of multiple phases. Crash recovery is optionally available: in case of failure the current state of the flow is saved for later inspection. A subclass of the basic flow class (`CheckpointFlow`) allows user-supplied checkpoint functions to be executed at the end of each phase, for example to save the internal structures of a node for later analysis. Flow objects are Python containers. Most of the builtin `list` methods are available. A `Flow` can be saved or copied using the corresponding `save` and `copy` methods.

Flow instantiation, training and execution

Suppose we have an input signal with an high number of dimensions, on which we would like to perform ICA. To make the problem affordable, we first need to reduce its dimensionality with PCA. Finally, we would like to find out the data that produces local maxima in the output on a new test set. This information could be used to characterize the input-output filters.

We start by generating some input signal at random (which makes the example useless, but it's just for illustration...). Generate 1000 observations of 20 independent source signals:

```
>>> inp = mdp.numx_rand.random((1000, 20))
```

Rescale x to have zero mean and unit variance:

```
>>> inp = (inp - mdp.numx.mean(inp, 0))/mdp.numx.std(inp, 0)
```

We reduce the variance of the last 15 components, so that they are going to be eliminated by PCA:

```
>>> inp[:,5:] /= 10.0
```

Mix the input signals linearly:

```
>>> x = mdp.utils.mult(inp,mdp.numx_rand.random((20, 20)))
```

`x` is now the training data for our simulation. In the same way we also create a test set `x_test`.

```
>>> inp_test = mdp.numx_rand.random((1000, 20))
>>> inp_test = (inp_test -
mdp.numx.mean(inp_test, 0))/mdp.numx.std(inp_test, 0)
>>> inp_test[:,5:] /= 10.0
>>> x_test = mdp.utils.mult(inp_test, mdp.numx_rand.random((20, 20)))
```

- We could now perform our analysis using only nodes, that's the lengthy way...

1. Perform PCA:

```
>>> pca = mdp.nodes.PCANode(output_dim=5)
>>> pca.train(x)
>>> out1 = pca.execute(x)
```

2. Perform ICA using CuBICA algorithm:

```
>>> ica = mdp.nodes.CuBICANode()
>>> ica.train(out1)
>>> out2 = ica.execute(out1)
```

3. Find the three largest local maxima in the output of the ICA node when applied to the test data, using a `HitParadeNode`:

```
>>> out1_test = pca.execute(x_test)
>>> out2_test = ica.execute(out1_test)
>>> hitnode = mdp.nodes.HitParadeNode(3)
>>> hitnode.train(out2_test)
>>> maxima, indices = hitnode.get_maxima()
```

- ... or we could use flows, which is the best way:

```
>>> flow = mdp.Flow([mdp.nodes.PCANode(output_dim=5), mdp.nodes.CuBICANode()])
>>> flow.train(x)
```

Now the training phase of PCA and ICA are completed. Next we append a `HitParadeNode` which we want to train on the test data:

```
>>> flow.append(mdp.nodes.HitParadeNode(3))
>>> flow.train(x_test)
>>> maxima, indices = flow[2].get_maxima()
```

Just to check that everything works properly, we can calculate covariance between the generated sources and the output (should be approximately 1):

```
>>> out = flow.execute(x)
>>> cov = mdp.numx.amax(abs(mdp.utils.cov2(inp[:, :5], out)))
>>> print cov
[ 0.98992083  0.99244511  0.99227319  0.99663185  0.9871812 ]
```

The `HitParadeNode` is an analysis node and as such does not interfere with the data flow.

Flow inversion

Flows can be inverted by calling their `inverse` method. In the case where the flow contains non-invertible nodes, trying to invert it would raise an exception. In this case, however, all nodes are invertible. We can reconstruct the mix by inverting the flow:

```
>>> rec = flow.inverse(out)
```

Calculate covariance between input mix and reconstructed mix: (should be approximately 1)

```
>>> cov = mdp.numx.amax(abs(mdp.utils.cov2(x/mdp.numx.std(x,axis=0),
...                                     rec/mdp.numx.std(rec,axis=0))))
>>> print cov
[ 0.99839606  0.99744461  0.99616208  0.99772863  0.99690947
  0.99864056  0.99734378  0.98722502  0.98118101  0.99407939
  0.99683096  0.99756988  0.99664384  0.99723419  0.9985529
  0.99829763  0.9982712  0.99721741  0.99682906  0.98858858]
```

Flows are container type objects

Flows are Python container type objects, very much like lists, i.e., you can loop through them:

```
>>> for node in flow:
...     print repr(node)
...
PCANode(input_dim=20, output_dim=5, dtype='float64')
CuBICANode(input_dim=5, output_dim=5, dtype='float64')
HitParadeNode(input_dim=5, output_dim=5, dtype='float64')
>>>
```

You can get slices, pop, insert, and append nodes like you would do with lists:

```
>>> len(flow)
3
>>> print flow[:2]
[PCANode, HitParadeNode]
>>> nodetoberemoved = flow.pop(-1)
>>> nodetoberemoved
HitParadeNode(input_dim=5, output_dim=5, dtype='float64')
>>> len(flow)
2
```

Finally, you can concatenate flows:

```
>>> dummyflow = flow[1:].copy()
>>> longflow = flow + dummyflow
>>> len(longflow)
3
```

The returned flow must always be consistent, i.e. input and output dimensions of successive nodes always have to match. If you try to create an inconsistent flow you'll get an error.

Crash recovery

If a node in a flow fails, you'll get a traceback that tells you which node has failed. You can also switch the crash recovery capability on. If something goes wrong you'll end up with a pickle dump of the flow, that can be later inspected.

To see how it works let's define a bogus node that always throws an `Exception` and put it into a flow:

```
>>> class BogusExceptNode(mdp.Node):
...     def train(self,x):
...         self.bogus_attr = 1
...         raise Exception, "Bogus Exception"
...     def execute(self,x):
...         raise Exception, "Bogus Exception"
...
>>> flow = mdp.Flow([BogusExceptNode()])
```

Switch on crash recovery:

```
>>> flow.set_crash_recovery(1)
```

Attempt to train the flow:

```
>>> flow.train(x)
Traceback (most recent call last):
  File "<stdin>", line 1, in ?
  [...]
mdp.linear_flows.FlowExceptionCR:
-----
! Exception in node #0 (BogusExceptNode):
Node Traceback:
Traceback (most recent call last):
  [...]
Exception: Bogus Exception
-----
A crash dump is available on: "/tmp/MDPcrash_LmISO_.pic"
```

You can give a file name to tell the flow where to save the dump:

```
>>> flow.set_crash_recovery('/home/myself/mydumps/MDPdumppic')
```

Iterators

Python allows user-defined classes to support iteration, as described in the [Python docs](#). A convenient implementation of the iterator protocol is provided by generators: see [this article](#) for an introduction, and the [official PEP](#) for a complete description.

Let us define two bogus node classes to be used as examples of nodes:

```
>>> class BogusNode(mdp.Node):
...     """This node does nothing."""
...     def _train(self, x):
...         pass
...
>>> class BogusNode2(mdp.Node):
...     """This node does nothing. But it's not trainable nor invertible.
...     """
...     def is_trainable(self): return False
...     def is_invertible(self): return False
...
>>>
```

This generator generates blocks input blocks to be used as training set. In this example one block is a 2-dimensional time-series. The first variable is [2,4,6,...,1000] and the second one [0,1,3,5,...,999]. All blocks are equal, this of course would not be the case in a real-life example.

In this example we use a progress bar to get progress information.

```
>>> def gen_data(blocks):
...     for i in mdp.utils.progressinfo(xrange(blocks)):
...         block_x = mdp.numx.atleast_2d(mdp.numx.arange(2,1001,2))
...         block_y = mdp.numx.atleast_2d(mdp.numx.arange(1,1001,2))
...         # put variables on columns and observations on rows
...         block = mdp.numx.transpose(mdp.numx.concatenate([block_x,block_y]))
...         yield block
...
>>>
```

Let's define a bogus flow consisting of 2 BogusNode:

```
>>> flow = mdp.Flow([BogusNode(),BogusNode()], verbose=1)
```

Train the first node with 5000 blocks and the second node with 3000 blocks. Note that the only allowed argument to `train` is a sequence (list or tuple) of iterators. In case you don't want or need to use incremental learning and want to do a one-shot training, you can use as argument to `train` a single array of data:

block-mode training

```
>>> flow.train([gen_data(5000),gen_data(3000)])
Training node #0 (BogusNode)
[=====100%=====>]

Training finished
Training node #1 (BogusNode)
[=====100%=====>]

Training finished
Close the training phase of the last node
```

one-shot training using one single set of data for both nodes

```
>>> flow = mdp.Flow([BogusNode(),BogusNode()])
>>> block_x = mdp.numx.atleast_2d(mdp.numx.arange(2,1001,2))
>>> block_y = mdp.numx.atleast_2d(mdp.numx.arange(1,1001,2))
>>> sin-
gle_block = mdp.numx.transpose(mdp.numx.concatenate([block_x,block_y]))
>>> flow.train(single_block)
```

If your flow contains non-trainable nodes, you must specify a `None` iterator for the non-trainable nodes:

```
>>> flow = mdp.Flow([BogusNode2(),BogusNode()], verbose=1)
>>> flow.train([None, gen_data(5000)])
Training node #0 (BogusNode2)
Training finished
Training node #1 (BogusNode)
[=====100%=====>]

Training finished
Close the training phase of the last node
```

You can use the one-shot training:

```
>>> flow = mdp.Flow([BogusNode2(),BogusNode()], verbose=1)
>>> flow.train(single_block)
Training node #0 (BogusNode2)
Training finished
Training node #1 (BogusNode)
Training finished
Close the training phase of the last node
```

Iterators can be used also for execution (and inversion):

```
>>> flow = mdp.Flow([BogusNode(),BogusNode()], verbose=1)
>>> flow.train([gen_data(1), gen_data(1)])
Training node #0 (BogusNode2)
```



```

Training finished
Training node #1 (IdentityNode)
[=====100%=====>]

```

```

Training finished
Close the training phase of the last node
>>> output = flow.execute(gen_data(1000))
[=====100%=====>]
>>> output = flow.inverse(gen_data(1000))
[=====100%=====>]

```

Execution and inversion can be done in one-shot mode also. Note that since training is finished you are not going to get a warning

```

>>> output = flow.execute(single_block)
>>> output = flow.inverse(single_block)

```

If a node requires multiple training phases (e.g., `GaussianClassifierNode`), Flow automatically takes care of reusing the iterator multiple times. In this case generators are not allowed, since they *expire* after yielding the last data block. If you try to restart them, they raise a `StopIteration` exception. General iterators, instead, can always be restarted. For example, you can loop over a list as many times as you need.

However, it is fairly easy to wrap a generator in a simple iterator if you need to:

```

>>> class SimpleIterator(object):
...     def __init__(self, blocks):
...         self.blocks = blocks
...     def __iter__(self):
...         # this is a generator
...         for i in range(self.blocks):
...             yield generate_some_data()
>>>

```

Note that if you use random numbers within the iterator, you usually would like to reset the random number generator to produce the same sequence every time:

```

>>> class RandomIterator(object):
...     def __init__(self):
...         self.state = None
...     def __iter__(self):
...         if self.state is None:
...             self.state = mdp.numx_rand.get_state()
...         else:
...             mdp.numx_rand.set_state(self.state)
...         for i in range(2):
...             yield mdp.numx_rand.random((1,4))
>>> iterator = RandomIterator()
>>> for x in iterator: print x
...
[[ 0.99586495  0.53463386  0.6306412  0.09679571]]
[[ 0.51117469  0.46647448  0.95089738  0.94837122]]
>>> for x in iterator: print x
...
[[ 0.99586495  0.53463386  0.6306412  0.09679571]]
[[ 0.51117469  0.46647448  0.95089738  0.94837122]]

```

Checkpoints

It can sometimes be useful to execute arbitrary functions at the end of the training or execution phase, for example to save the internal structures of a node for later analysis. This can easily be done by defining a `CheckpointFlow`. As an example imagine the following situation: you want to perform Principal Component Analysis (PCA) on your data to reduce the dimensionality. After this you want to expand the signals into a nonlinear space and then perform Slow Feature Analysis to extract slowly varying signals. As the expansion will increase the number of components, you don't want to run out of memory, but at the same time you want to keep as much information as possible after the dimensionality reduction. You could do that by specifying the percentage of the total input variance that has to be conserved in the dimensionality reduction. As the number of output components of the PCA node now can become as large as the that of the input components, you want to check, after training the PCA node, that this number is below a certain threshold. If this is not the case you want to abort the execution and maybe start again requesting less variance to be kept.

Let start defining a generator to be used through the whole example:

```
>>> def gen_data(blocks,dims):
...     mat = mdp.numx_rand.random((dims,dims))-0.5
...     for i in xrange(blocks):
...         # put variables on columns and observations on rows
...         block = mdp.utils.mult(mdp.numx_rand.random((1000,dims)), mat)
...         yield block
...
>>>
```

Define a `PCANode` which reduces dimensionality of the input, a `PolynomialExpansionNode` to expand the signals in the space of polynomials of degree 2 and a `SFANode` to perform SFA:

```
>>> pca = mdp.nodes.PCANode(output_dim=0.9)
>>> exp = mdp.nodes.PolynomialExpansionNode(2)
>>> sfa = mdp.nodes.SFANode()
```

As you see we have set the output dimension of the `PCANode` to be 0.9. This means that we want to keep at least 90% of the variance of the original signal. We define a `PCADimensionExceededException` that has to be thrown when the number of output components exceeds a certain threshold:

```
>>> class PCADimensionExceededException(Exception):
...     """Exception base class for PCA exceeded dimensions case."""
...     pass
...
>>>
```

Then, write a `CheckpointFunction` that checks the number of output dimensions of the `PCANode` and aborts if this number is larger than `max_dim`:

```
>>> class CheckPCA(mdp.CheckpointFunction):
...     def __init__(self,max_dim):
...         self.max_dim = max_dim
...     def __call__(self,node):
...         node.stop_training()
...         act_dim = node.get_output_dim()
...         if act_dim > self.max_dim:
...             errstr = 'PCA output dimensions exceeded maximum '+\
...                 '(%d > %d)'%(act_dim,self.max_dim)
...             raise PCADimensionExceededException, errstr
...         else:
```

```

...             print 'PCA output dimensions = %d'%(act_dim)
...
>>>

```

Define the CheckpointFlow:

```
>>> flow = mdp.CheckpointFlow([pca, exp, sfa])
```

To train it we have to supply 3 generators and 3 checkpoint functions:

```

>>> flow.train([gen_data(10, 50), None, gen_data(10, 50)],
...             [CheckPCA(10), None, None])
Traceback (most recent call last):
  File "<stdin>", line 2, in ?
  [...]
__main__.PCADimensionExceededException: PCA output dimensions exceeded maxi-
mum (25 > 10)

```

The training fails with a PCADimensionExceededException. If we only had 12 input dimensions instead of 50 we would have passed the checkpoint:

```

>>> flow[0] = mdp.nodes.PCANode(output_dim=0.9)
>>> flow.train([gen_data(10, 12), None, gen_data(10, 12)],
...             [CheckPCA(10), None, None])
PCA output dimensions = 6

```

We could use the built-in CheckpointSaveFunction to save the SFANode and analyze the results later :

```

>>> pca = mdp.nodes.PCANode(output_dim=0.9)
>>> exp = mdp.nodes.PolynomialExpansionNode(2)
>>> sfa = mdp.nodes.SFANode()
>>> flow = mdp.CheckpointFlow([pca, exp, sfa])
>>> flow.train([gen_data(10, 12), None, gen_data(10, 12)],
...             [CheckPCA(10),
...             None,
...             mdp.CheckpointSaveFunction('dummy.pic',
...                                         stop_training = 1,
...                                         protocol = 0)])
...
PCA output dimensions = 7

```

We can now reload and analyze the SFANode:

```

>>> fl = file('dummy.pic')
>>> import cPickle
>>> sfa_reloaded = cPickle.load(fl)
>>> sfa_reloaded
SFANode(input_dim=35, output_dim=35, dtype='d')

```

Don't forget to clean the rubbish:

```

>>> fl.close()
>>> import os
>>> os.remove('dummy.pic')

```

A real life example (Logistic maps)

We show an application of Slow Feature Analysis to the analysis of non-stationary time series. We consider a chaotic time series generated by the logistic map based on the logistic equation (a demographic model of the population biomass of species in the presence of limiting factors such as food supply or disease), and extract the slowly varying parameter that is hidden behind the time series. This example reproduces some of the results reported in: Laurenz Wiskott, *Estimating Driving Forces of Nonstationary Time Series with Slow Feature Analysis*. arXiv.org e-Print archive, <http://arxiv.org/abs/cond-mat/0312317>

Generate the slowly varying driving force, a combination of three sine waves (freqs: 5, 11, 13 Hz), and define a function to generate the logistic map

```
>>> p2 = mdp.numx.pi*2
>>> t = mdp.numx.linspace(0,1,10000,endpoint=0) # time axis 1s, sampler-
ate 10KHz
>>> dforce = mdp.numx.sin(p2*5*t) + mdp.numx.sin(p2*11*t) + mdp.numx.sin(p2*13*t)
>>> def logistic_map(x,r):
...     return r*x*(1-x)
...
>>>
```

Note that we define **series** to be a two-dimensional array. Inputs to MDP must be two-dimensional arrays with variables on columns and observations on rows. In this case we have only one variable:

```
>>> series = mdp.numx.zeros((10000,1),'d')
```


Fix the initial condition:

```
>>> series[0] = 0.6
```

Generate the time-series using the logistic equation the driving force modifies the logistic equation parameter **r**:

```
>>> for i in range(1,10000):
...     series[i] = logistic_map(series[i-1],3.6+0.13*dforce[i])
...
>>>
```

If you have a plotting package **series** should look like this:



Define a flow to perform SFA in the space of polynomials of degree 3. We need a node that embeds the time-series in a 10 dimensional space, where different variables correspond to time-delayed copies of the original time-series: the `TimeFramesNode(10)`. Then we need a node that expands the new signal in the space of polynomials of degree 3: the `PolynomialExpansionNode(3)`. Finally we perform SFA onto the expanded signal and keep the slowest feature: `SFANode(output_dim=1)`. We also measure the *slowness* of the input time-series and of the slow feature obtained by SFA. Therefore we put at the beginning and at the end of the sequence an *analysis node* that computes the *eta-value* (a measure of slowness) of its input (see docs for the definition of eta-value): the `EtaComputerNode()`:

```
>>> sequence = [mdp.nodes.EtaComputerNode(),
...             mdp.nodes.TimeFramesNode(10),
...             mdp.nodes.PolynomialExpansionNode(3),
...             mdp.nodes.SFANode(output_dim=1),
...             mdp.nodes.EtaComputerNode()]
>>>
>>> flow = mdp.Flow(sequence, verbose=1)
```

Since the time-series is short enough to be kept in memory we don't need to define generators and we can feed the flow directly with the whole signal:

```
>>> flow.train(series)
```

Since the second and the third nodes are not trainable we are going to get two warnings (**Training Interrupted**). We can safely ignore them. Execute the flow to get the slow feature

```
>>> slow = flow(series)
```

The slow feature should match the driving force up to a scaling factor, a constant offset and the sign. To allow a comparison we rescale the driving force to have zero mean and unit variance:

```
>>> resc_dforce = (dforce - mdp.numx.mean(dforce,0))/mdp.numx.std(dforce,0)
```


Print covariance between the rescaled driving force and the slow feature. Note that embedding the time-series with 10 time frames leads to a time-series with 9 observations less:

```
>>> mdp.utils.cov2(resc_dforce[:-9],slow)
0.99992501533859179
```

Print the *eta-values* of the chaotic time-series and of the slow feature

```
>>> print 'Eta value (time-series): ', flow[0].get_eta(t=10000)
Eta value (time-series): [ 3002.53380245]
>>> print 'Eta value (slow feature): ', flow[-1].get_eta(t=9996)
Eta value (slow feature): [ 10.2185087]
```

If you have a plotting package you could plot `resc_dforce` together with `slow` and see that they match perfectly:



Another real life example (Growing neural gas)

We generate uniformly distributed random data points confined on different 2-D geometrical objects. The Growing Neural Gas Node builds a graph with the same topological structure.

Fix the random seed to obtain reproducible results:

```
>>> mdp.numx_rand.seed(1266090063)
```

Some functions to generate uniform probability distributions on different geometrical objects:

```
>>> def uniform(min_, max_, dims):
...     """Return a random number between min_ and max_ ."""
...     return mdp.numx_rand.random(dims)*(max_-min_)+min_
...
>>> def circumference_distr(center, radius, n):
...     """Return n random points uniformly distributed on a circumfer-
...     ence."""
...     phi = uniform(0, 2*mdp.numx.pi, (n,1))
...     x = radius*mdp.numx.cos(phi)+center[0]
...     y = radius*mdp.numx.sin(phi)+center[1]
```

```

...     return mdp.numx.concatenate((x,y), axis=1)
...
>>> def circle_distr(center, radius, n):
...     """Return n random points uniformly distributed on a circle."""
...     phi = uniform(0, 2*mdp.numx.pi, (n,1))
...     sqrt_r = mdp.numx.sqrt(uniform(0, radius*radius, (n,1)))
...     x = sqrt_r*mdp.numx.cos(phi)+center[0]
...     y = sqrt_r*mdp.numx.sin(phi)+center[1]
...     return mdp.numx.concatenate((x,y), axis=1)
...
>>> def rectangle_distr(center, w, h, n):
...     """Return n random points uniformly distributed on a rectangle."""
...     x = uniform(-w/2., w/2., (n,1))+center[0]
...     y = uniform(-h/2., h/2., (n,1))+center[1]
...     return mdp.numx.concatenate((x,y), axis=1)
...
>>> N = 2000

```

Explicitly collect random points from some distributions:

- Circumferences:

```

>>> cf1 = circumference_distr([6,-0.5], 2, N)
>>> cf2 = circumference_distr([3,-2], 0.3, N)

```

- Circles:

```

>>> cl1 = circle_distr([-5,3], 0.5, N/2)
>>> cl2 = circle_distr([3.5,2.5], 0.7, N)

```

- Rectangles:

```

>>> r1 = rectangle_distr([-1.5,0], 1, 4, N)
>>> r2 = rectangle_distr([+1.5,0], 1, 4, N)
>>> r3 = rectangle_distr([0,+1.5], 2, 1, N/2)
>>> r4 = rectangle_distr([0,-1.5], 2, 1, N/2)

```


Shuffle the points to make the statistics stationary

```

>>> x = mdp.numx.concatenate([cf1, cf2, cl1, cl2, r1,r2,r3,r4], axis=0)
>>> x = mdp.numx.take(x,mdp.numx_rand.permutation(x.shape[0]), axis=0)

```


If you have a plotting package x should look like this:



Create a `GrowingNeuralGasNode` and train it:

```
>>> gng = mdp.nodes.GrowingNeuralGasNode(max_nodes=75)
```

The initial distribution of nodes is randomly chosen:




The training is performed in small chunks in order to visualize the evolution of the graph:

```
>>> STEP = 500
>>> for i in range(0,x.shape[0],STEP):
...     gng.train(x[i:i+STEP])
...     # [...] plotting instructions
...
>>> gng.stop_training()
```

See [here](#) the animation of training.

Visualizing the neural gas network, we'll see that it is adapted to the topological structure of the data distribution:



Calculate the number of connected components:

```
>>> n_obj = len(gng.graph.connected_components())
>>> print n_obj
5
```

Node List

Here is the complete list of implemented nodes. Refer to the [API](#) for the full documentation and interface description.

CuBICANode Perform Independent Component Analysis using the CuBICA algorithm. Reference: Blaschke, T. and Wiskott, L. (2003). *CuBICA: Independent Component Analysis by Simultaneous Third- and Fourth-Order Cumulant Diagonalization*. IEEE Transactions on Signal Processing, 52(5), pp. 1250-1256. More information about ICA can be found among others in Hyvarinen A., Karhunen J., Oja E. (2001). *Independent Component Analysis*, Wiley.

EtaComputerNode Compute the eta values of the normalized training data. The delta value of a signal is a measure of its temporal variation, and is defined as the mean of the derivative squared, i.e. $\text{delta}(x) = \text{mean}(\text{dx}/\text{dt}(t)^2)$. $\text{delta}(x)$ is zero if 'x' is a constant signal, and increases if the temporal variation of the signal is bigger. The eta value is a more intuitive measure of temporal variation, defined as $\text{eta}(x) = T/(2\pi) * \text{sqrt}(\text{delta}(x))$. If 'x' is a signal of length 'T' which consists of a sine function that accomplishes exactly 'N' oscillations, then $\text{eta}(x) = N$. Reference: Wiskott, L. and Sejnowski, T.J. (2002). *Slow Feature Analysis: Unsupervised Learning of Invariances*, Neural Computation, 14(4):715-770.

FANode Perform Factor Analysis. The current implementation should be most efficient for long data sets: the sufficient statistics are collected in the training phase, and all EM-cycles are performed at its end. More information about Factor Analysis can be found in [Max Welling's classnotes](#) in the chapter "Linear Models".

FastICANode Perform Independent Component Analysis using the FastICA algorithm. Reference: Aapo Hyvarinen (1999). *Fast and Robust Fixed-Point Algorithms for Independent Component Analysis*, IEEE Transactions on Neural Networks, 10(3):626-634. More information about ICA can be found among others in Hyvarinen A., Karhunen J., Oja E. (2001). *Independent Component Analysis*, Wiley.

FDANode Perform a (generalized) Fisher Discriminant Analysis of its input. It is a supervised node that implements FDA using a generalized eigenvalue approach. More information on Fisher Discriminant Analysis can be found for example in C. Bishop, *Neural Networks for Pattern Recognition*, Oxford Press, pp. 105-112.

GaussianClassifierNode Perform a supervised Gaussian classification. Given a set of labelled data, the node fits a gaussian distribution to each class.

GrowingNeuralGasNode Learn the topological structure of the input data by building a corresponding graph approximation. More information about the Growing Neural Gas algorithm can be found in B. Fritzke, *A Growing Neural Gas Network Learns Topologies*, in G. Tesauero, D. S. Touretzky, and T. K. Leen (editors), *Advances in Neural Information Processing Systems 7*, pages 625-632. MIT Press, Cambridge MA, 1995.

HitParadeNode Collect the first 'n' local maxima and minima of the training signal which are separated by a minimum gap 'd'.

ISFANode Perform Independent Slow Feature Analysis on the input data. More information about ISFA can be found in: Blaschke, T. , Zito, T., and Wiskott, L. *Independent Slow Feature Analysis and Nonlinear Blind Source Separation*. Neural Computation 19(4):994-1021 (2007).

NoiseNode Inject multiplicative or additive noise into the input data.

PCANode Filter the input data through the most significant of its principal components. More information about Principal Component Analysis, a.k.a. discrete Karhunen-Loeve transform can be found among others in I.T. Jolliffe, *Principal Component Analysis*, Springer-Verlag (1986).

PolynomialExpansionNode Perform expansion in a polynomial space.

QuadraticExpansionNode Perform expansion in the space formed by all linear and quadratic monomials

SFANode Extract the slowly varying components from the input data. More information about Slow Feature Analysis can be found in Wiskott, L. and Sejnowski, T.J., *Slow Feature Analysis: Unsupervised Learning of Invariances*, Neural Computation, 14(4):715-770 (2002).

SFA2Node Get an input signal, expand it in the space of inhomogeneous polynomials of degree 2 and extract its slowly varying components. The `get_quadratic_form` method returns the input-output function of one of the learned unit as a `mdp.utils.QuadraticForm` object. More information about Slow Feature Analysis can be found in Wiskott, L. and Sejnowski, T.J., *Slow Feature Analysis: Unsupervised Learning of Invariances*, Neural Computation, 14(4):715-770 (2002).

TimeFramesNode Copy delayed version of the input signal on the space dimensions.

For example, for `time_frames=3` and `gap=2`:

[X(1) Y(1)		[X(1) Y(1) X(3) Y(3) X(5) Y(5)
X(2) Y(2)		X(2) Y(2) X(4) Y(4) X(6) Y(6)
X(3) Y(3)	-->	X(3) Y(3) X(5) Y(5) X(7) Y(7)
X(4) Y(4)		X(4) Y(4) X(6) Y(6) X(8) Y(8)
X(5) Y(5)	
X(6) Y(6)		
X(7) Y(7)		
X(8) Y(8)		
... ..]		

WhiteningNode 'Whiten' the input data by filtering it through the most significant of its principal components. All output signals have zero mean, unit variance and are decorrelated.

Contributed Nodes

Here is a list of nodes that have been contributed by MDP users:

JADENode Perform Independent Component Analysis using the JADE algorithm. References: Cardoso, J.-F, and Souloumiac, A. *Blind beamforming for non Gaussian signals*. Radar and Signal Processing, IEE Proceedings F, 140(6): 362-370 (1993), and Cardoso, J.-F. *High-order contrasts for independent component analysis*. Neural Computation, 11(1): 157-192 (1999). More information about ICA can be found among others in Hyvarinen A., Karhunen J., Oja E. (2001). *Independent Component Analysis*, Wiley.

Code contributed by: Gabriel Beckers

NIPALSNode Perform Principal Component Analysis using the NIPALS algorithm. Reference for NIPALS (Nonlinear Iterative Partial Least Squares): Wold, H. *Nonlinear estimation by iterative least squares procedures*. in David, F. (Editor), Research Papers in Statistics, Wiley, New York, pp 411-444 (1966).

More information about Principal Component Analysis, a.k.a. discrete Karhunen-Loeve transform can be found among others in I.T. Jolliffe, *Principal Component Analysis*, Springer-Verlag (1986).

Code contributed by: Michael Schmuker, Susanne Lezius, and Farzad Farkhooi (2008).

Didn't you find what you were looking for?

If you want to contribute some code or a new algorithm, please do not hesitate to submit it!

Additional utilities

MDP offers some additional utilities of general interest in the `mdp.utils` module. Refer to the [API](#) for the full documentation and interface description.

CovarianceMatrix This class stores an empirical covariance matrix that can be updated incrementally. A call to the `fix` method returns the current state of the covariance matrix, the average and the number of observations, and resets the internal data.

Note that the internal sum is a standard `__add__` operation. We are not using any of the fancy sum algorithms to avoid round off errors when adding many numbers. If you want to contribute a `CovarianceMatrix` class that uses such algorithms we would be happy to include it in MDP. For a start see the [Python recipe](#) by Raymond Hettinger. For a review about floating point arithmetic and its pitfalls see this [interesting article](#).

DelayCovarianceMatrix This class stores an empirical covariance matrix between the signal and time delayed signal that can be updated incrementally.

MultipleCovarianceMatrices Container class for multiple covariance matrices to easily execute operations on all matrices at the same time.

dig_node(node) Crawl recursively an MDP Node looking for arrays. Return (dictionary, string), where the dictionary is: { `attribute_name`: (`size_in_bytes`, `array_reference`)} and string is a nice string representation of it.

get_node_size(node) Get 'node' total byte-size using `cPickle` with `protocol=2`. (The byte-size is related the memory needed by the node).

progressinfo(sequence, length, style, custom) A fully configurable text-mode progress info box. To get a progress info box for your loops use it like this:

```
>>> for i in progressinfo(sequence):
...     do_something(i)
```

You can also use it with generators, files or any other iterable object, but in this case you have to specify the total length of the sequence:

```
>>> for line in progressinfo(open_file, nlines):
...     do_something(line)
```

A few examples of the available layouts:

```
[=====73%=====>.....]

Progress: 67%[=====>

23% [02:01:28] - [00:12:37]
```

QuadraticForm Define an inhomogeneous quadratic form as $1/2 \mathbf{x}'\mathbf{H}\mathbf{x} + \mathbf{f}'\mathbf{x} + c$. This class implements the quadratic form analysis methods presented in: Berkes, P. and Wiskott, L. On the analysis and interpretation of inhomogeneous quadratic forms as receptive fields. *Neural Computation*, 18(8): 1868-1895. (2006).

refcast(array, dtype) Cast the array to 'dtype' only if necessary, otherwise return a reference.

rotate(mat, angle, columns, units) Rotate in-place a NxM data matrix in the plane defined by the 'columns' when observation are stored on rows. Observations are rotated counterclockwise. This corresponds to the following matrix-multiplication for each data-point (unchanged elements omitted):

$$\begin{bmatrix} \cos(\text{angle}) & -\sin(\text{angle}) \\ \sin(\text{angle}) & \cos(\text{angle}) \end{bmatrix} * \begin{bmatrix} x_i \\ x_j \end{bmatrix}$$

random_rot(dim, dtype) Return a random rotation matrix, drawn from the Haar distribution (the only uniform distribution on $SO(n)$). The algorithm is described in the paper Stewart, G.W., *The efficient generation of random orthogonal matrices with an application to condition estimators*, SIAM Journal on Numerical Analysis, 17(3), pp. 403-409, 1980. For more information see this [Wikipedia entry](#).

symrand(dim_or_eigv, dtype) Return a random symmetric (Hermitian) matrix with eigenvalues uniformly distributed on (0,1].


Graph module

MDP contains `mdp.graph`, a lightweight package to handle directed graphs.

Graph Represent a directed graph. This class contains several methods to create graph structures and manipulate them, among which

- **add_tree: Add a tree to the graph.** The tree is specified with a nested list of tuple, in a LISP-like notation. The values specified in the list become the values of the single nodes. Return an equivalent nested list with the nodes instead of the values.
Example:

```
>>> a=b=c=d=e=None
>>> g.add_tree( (a, b, (c, d ,e)) )
# corresponds to this tree structure, with all node val-
ues set to None:
```



- `topological_sort`: Perform a topological sort of the nodes.
- `dfs, undirected_dfs`: Perform Depth First sort.
- `bfs, undirected_bfs`: Perform Breadth First sort.
- `connected_components`: **Return a list of lists containing** the nodes of all connected components of the graph.
- `is_weakly_connected`: Return True if the graph is weakly connected.

GraphEdge Represent a graph edge and all information attached to it.

GraphNode Represent a graph node and all information attached to it.

recursive_map(func, seq) Apply a function recursively on a sequence and all subsequences.

recursive_reduce(func, seq, *argv) Apply `reduce(func, seq)` recursively to a sequence and all its subsequences.

To Do

In this section we want to give you an overview about our plans for the development of MDP:

- Add more data processing algorithms.
- Extend the linear flows to handle general acyclic graphs of nodes.
- Actual use of the graph structure will be possible only in presence of an easy and intuitive GUI :)
- Wait for a good guy who wants to contribute a `CovarianceMatrix` class that uses some of the fancy sum algorithms to avoid round off errors when adding many numbers.

Contributors

In this final section we want to thank all users who have contributed code to the MDP project. Strictly in alphabetical order:

- [Gabriel Beckers](#)
- [Farzad Farkhooi](#)
- Susanne Lezius
- [Michael Schmuker](#)

Generated on: 2008-03-13 16:00 UTC. Generated by [Docutils](#) from [reStructuredText](#) source.